

Project Milestone 1-3 Report

Caiting Wu (cwu72)
Hua Chen (huachen4)
Zekun Hu (zekunhu2)

Milestone 1

1. A List of all Kernels that Collectively Consume More than 90% of the Program Time:

| Time(%) | Time(ms) | Name |
|---------|----------|--|
| 34.00 | 118.46 | void fermiPlusCgemmLDS128_batched<...> |
| 26.94 | 93.879 | void cudnn::detail::implicit_convolve_sgemm<...> |
| 12.65 | 44.060 | void fft2d_c2r_32x32<...> |
| 8.19 | 28.540 | sgemm_sm35_ldg_tn_128x8x256x16x32 |
| 6.65 | 23.153 | [CUDA memcpy HtoD] |
| 4.06 | 14.157 | void cudnn::detail::activation_fw_4d_kernel<...> |
| 3.81 | 13.289 | void cudnn::detail::pooling_fw_4d_kernel<...> |
| 1.71 | 5.9454 | void fft2d_r2c_32x32<...> |
| 1.16 | 4.0542 | sgemm_sm35_ldg_tn_64x16x128x8x32 |

2. A list of all CUDA API calls that collectively consume more than 90% of the program time:

| Time(%) | Time(ms) | Name |
|---------|----------|---------------------------|
| 43.48 | 1.92 | cudaStreamCreateWithFlags |
| 27.11 | 1.20 | cudaFree |
| 20.71 | 916.54 | cudaMemGetInfo |
| 7.31 | 323.45 | cudaStreamSynchronize |
| 1.08 | 47.75 | cudaMemcpy2DAsync |
| 0.16 | 7.15 | cudaMalloc |
| 0.03 | 1.37 | cuDeviceTotalMem |

3. Explain the difference between kernels and API calls

Kernels are function launched to be executed on the device while APIs are used to communicate between different software components.

4. Show Output of rai Running MXNet on the CPU:

```
Successfully installed mxnet
* Running /usr/bin/time python ml.1.py
Loading fashion-mnist data...
done
Loading model...
done
New Inference
EvalMetric: {'accuracy': 0.8444}
13.16user 12.04system 0:11.76elapsed 214%CPU (0avgtext+0avgdata 2830016maxresident)k
0inputs+2624outputs (0major+36851minor)pagefaults 0swaps
```

Figure: the snapshot of the result we have by running MXNet on the CPU

5. Program Run-Time (CPU):

From the result shown in the PowerShell, we can tell the run-time on the CPU is 11.76 seconds.

6. Show Output of rai Running MXNet on the GPU:

```
Successfully installed mxnet
* Running /usr/bin/time python ml.2.py
Loading fashion-mnist data...
done
Loading model...
[04:54:48] src/operator/././cudnn_algoreg-inl.h:112: Running performance tests to find
o disable)
done
New Inference
EvalMetric: {'accuracy': 0.8444}
2.22user 1.09system 0:02.81elapsed 117%CPU (0avgtext+0avgdata 1139152maxresident)k
0inputs+3136outputs (0major+159479minor)pagefa
ults 0swaps
```

Figure: the snapshot of the result we have by running MXNet on the GPU

7. Program Run-Time (GPU):

From the result shown in the PowerShell, we can tell the run-time on the GPU is 2.81 seconds.

Milestone 2

1. List whole program execution time:

a). 10000 images (default): 30.10 s

```
Successfully installed mxnet
* Running /usr/bin/time python m2.1.py
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 6.607474
Op Time: 19.537141
Correctness: 0.8451 Model: ece408
30.64user 1.48system 0:30.10elapsed 106%CPU (0avgtext+0avgdata 2821096maxresident)k
0inputs+2624outputs (0major+37057minor)pagefaults 0swaps
```

Figure: the snapshot of the result we have by running ConvNet on the CPU with 10000 images

b). 100 images: 1.15 s

```
Successfully installed mxnet
* Running /usr/bin/time python m2.1.py 100
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 0.065591
Op Time: 0.194542
Correctness: 0.88 Model: ece408
1.20user 0.58system 0:01.15elapsed 155%CPU (0avgtext+0avgdata 187088maxresident)k
0inputs+2624outputs (0major+33639minor)pagefaults 0swaps
```

Figure: the snapshot of the result we have by running ConvNet on the CPU with 100 images

c). 10 images (default): 0.87 s

```
Successfully installed mxnet
* Running /usr/bin/time python m2.1.py 10
Loading fashion-mnist data...
done
Loading model...
done
New Inference
Op Time: 0.006572
Op Time: 0.019520
Correctness: 1.0 Model: ece408
0.86user 0.51system 0:00.87elapsed 156%CPU (0avgtext+0avgdata 170392maxresident)k
0inputs+2624outputs (0major+31144minor)pag
efaults 0swaps
```

Figure: the snapshot of the result we have by running ConvNet on the CPU with 10 images

2. List Op Times:

a). 10000 images (default):

| | |
|-----------------------|-------------|
| First Layer Op Time: | 6.60747 s |
| Second Layer Op Time: | 19.537141 s |

b). 100 images:

| | |
|-----------------------|------------|
| First Layer Op Time: | 0.065591 s |
| Second Layer Op Time: | 0.194542 s |

c). 10 images:

| | |
|-----------------------|------------|
| First Layer Op Time: | 0.006572 s |
| Second Layer Op Time: | 0.019520 s |

Milestone 3

1. : nvprof profiling the execution (10000 default data size):

```
* Running nvprof python m3.1.py
Loading fashion-mnist data...
done
Loading model...
==314== NVPROF is profiling process 314, command: python m3.1.py
done
```

New Inference

Op Time: 0.355486

Op Time: 0.652564

Correctness: 0.8451 **Model:** ece408

==314== Profiling application: python m3.1.py

==314== **Profiling result:**

| Time (%) | Time | Calls | Name |
|----------|-----------|-------|---|
| 89.76 | 1.00787 s | 2 | mxnet::op::forward_kernel |
| 2.61 | 29.329 ms | 1 | sgemm_sm35_ldg_tn_128x8x256x16x32 |
| 2.60 | 29.224 ms | 14 | [CUDA memcpy HtoD] |
| 1.96 | 22.008 ms | 2 | mshadow::cuda::MapPlanLargeKernel<...> |
| 1.26 | 14.158 ms | 2 | cudnn::detail::activation_fw_4d_kernel<...> |
| 1.19 | 13.310 ms | 1 | cudnn::detail::pooling_fw_4d_kernel<...> |
| 0.36 | 4.0730 ms | 1 | sgemm_sm35_ldg_tn_64x16x128x8x32 |
| 0.11 | 1.2805 ms | 1 | mshadow::cuda::MapPlanLargeKernel<...> |
| 0.10 | 1.1112 ms | 1 | mshadow::cuda::SoftmaxKernel<...> |
| 0.02 | 177.44 us | 13 | mshadow::cuda::MapPlanKernel<...> |
| 0.01 | 146.82 us | 2 | mshadow::cuda::MapPlanKernel<...> |
| 0.01 | 129.60 us | 1 | sgemm_sm35_ldg_tn_32x16x64x8x16 |
| <0.01 | 23.264 us | 1 | mshadow::cuda::MapPlanKernel<...> |
| <0.01 | 9.6320 us | 1 | [CUDA memcpy DtoH] |

==314== API calls:

| Time(%) | Time | Call | Avg | Min | Max | Name |
|----------------------------------|----------|------|----------|----------|-----------|------------------------------|
| 37.07% | 2.01992s | 18 | 112.22ms | 18.700us | 1.00956s | cudaStreamCreateWithFlags |
| 23.45% | 1.27787s | 10 | 127.79ms | 1.0610us | 359.44ms | cudaFree |
| 18.90% | 1.02999s | 6 | 171.66ms | 16.423us | 652.50ms | cudaDeviceSynchronize |
| 18.12% | 987.04ms | 27 | 36.557ms | 249.68us | 978.97ms | cudaMemGetInfo |
| 1.15% | 62.495ms | 29 | 2.1550ms | 5.9750us | 32.128ms | cudaStreamSynchronize |
| 1.09% | 59.196ms | 9 | 6.5774ms | 10.270us | 28.498ms | cudaMemcpy2DAsync |
| 0.12% | 6.6691ms | 45 | 148.20us | 9.9270us | 950.38us | cudaMalloc |
| 0.03% | 1.4217ms | 4 | 355.43us | 336.63us | 403.41us | cuDeviceTotalMem |
| 0.02% | 1.0917ms | 352 | 3.1010us | 512ns | 88.433us | cuDeviceGetAttribute |
| 0.02% | 827.59us | 114 | 7.2590us | 897ns | 305.72us | cudaEventCreateWithFlags |
| 0.01% | 681.04us | 28 | 24.322us | 11.212us | 67.046us | cudaLaunch |
| 0.01% | 407.12us | 6 | 67.853us | 28.668us | 137.98us | cudaMemcpy |
| 0.01% | 287.64us | 4 | 71.909us | 36.702us | 160.69us | cudaStreamCreate |
| 0.00% | 134.58us | 168 | 801ns | 522ns | 13.878us | cudaSetupArgument |
| 0.00% | 120.73us | 4 | 30.182us | 18.224us | 36.266us | cuDeviceGetName |
| 0.00% | 96.465us | 104 | 927ns | 627ns | 2.3410us | cudaDeviceGetAttribute |
| 0.00% | 89.771us | 34 | 2.6400us | 925ns | 7.5090us | cudaSetDevice |
| 0.00% | 45.847us | 2 | 22.923us | 21.098u | 24.749 us | cudaStreamCreateWithPriority |
| 0.00% | 45.809us | 28 | 1.6360us | 779ns | 4.3910us | cudaConfigureCall |
| 0.00% | 30.516us | 10 | 3.0510us | 1.6350us | 7.0540us | cudaGetDevice |
| 0.00% | 17.136us | 20 | 856ns | 583ns | 1.2970us | cudaPeekAtLastError |
| 0.00% | 7.0090us | 6 | 1.1680us | 550ns | 2.5230us | cuDeviceGetCount |
| 0.00% | 6.0670us | 2 | 3.0330us | 2.5290us | 3.5380us | cudaStreamWaitEvent |
| 0.00% | 5.6780us | 6 | 946ns | 604ns | 1.2170us | cuDeviceGet |
| 0.00% | 5.3790us | 2 | 2.6890us | 1.8120us | 3.5670us | cudaEventRecord |
| 0.00% | 5.2660us | 1 | 5.2660us | 5.2660us | 5.2660us | cudaStreamGetPriority |
| 0.00% | 4.7120us | 5 | 942ns | 726ns | 1.2410us | cudaGetLastError |
| 0.00% | 4.1180us | 3 | 1.3720us | 1.2970us | 1.4390us | cuInit |
| 0.00% | 3.8380us | 2 | 1.9190us | 1.5640us | 2.2740us | |
| cudaDeviceGetStreamPriorityRange | | | | | | |
| 0.00% | 3.5850us | 3 | 1.1950us | 907ns | 1.3690us | cuDriverGetVersion |
| 0.00% | 1.7840us | 1 | 1.7840us | 1.7840us | 1.7840us | cudaGetDeviceCount |

2. nvprof profiling the execution (100 data size):

* Running nvprof python m3.1.py 100

Loading fashion-mnist data...

done

Loading model...

==311== NVPROF is profiling process 311, command: python m3.1.py 100

done

New Inference

Op Time: 0.003272

Op Time: 0.006654

Correctness: 0.88 Model: ece408

==311== Profiling application: python m3.1.py 100

==311== Profiling result:

| Time (%) | Time | Calls | Name |
|----------|-----------|-------|--|
| 78.67 | 9.8344 ms | 2 | mxnet::op::forward_kernel |
| 9.15 | 1.1435 ms | 14 | [CUDA memcpy HtoD] |
| 5.45 | 680.92 us | 1 | sgemm_largek_lds64 |
| 2.02 | 252.67 us | 2 | void mshadow::cuda::MapPlanKernel<...> |
| 1.58 | 197.18 us | 14 | void mshadow::cuda::MapPlanKernel<...> |
| 1.30 | 162.62 us | 2 | void cudnn::detail::activation_fw_4d_kernel<...> |
| 1.12 | 140.51 us | 1 | void cudnn::detail::pooling_fw_4d_kernel<...> |
| 0.33 | 41.248 us | 1 | sgemm_sm35_ldg_tn_32x16x64x8x16 |
| 0.15 | 18.239 us | 1 | void mshadow::cuda::SoftmaxKernel<...> |
| 0.09 | 11.808 us | 1 | void mshadow::cuda::MapPlanKernel<...> |
| 0.06 | 7.0720 us | 2 | void mshadow::cuda::MapPlanKernel<...> |
| 0.04 | 4.5760 us | 1 | [CUDA memcpy DtoH] |
| 0.03 | 3.9680 us | 1 | void scal_kernel<...> |
| 0.01 | 1.4080 us | 1 | [CUDA memset] |

==311== API calls:

| Time(%) | Time | Call | Avg | Min | Max | Name |
|---------|----------|------|----------|----------|----------|---------------------------|
| 47.00% | 1.99860s | 16 | 124.91ms | 18.774us | 999.00ms | cudaStreamCreateWithFlags |
| 29.76% | 1.26567s | 10 | 126.57ms | 1.3710us | 357.94ms | cudaFree |
| 22.72% | 966.04ms | 27 | 35.779ms | 337.00us | 957.04ms | cudaMemGetInfo |
| 0.24% | 10.142ms | 6 | 1.6903ms | 4.9720us | 6.6253ms | cudaDeviceSynchronize |
| 0.07% | 2.8118ms | 9 | 312.43us | 10.485us | 783.21us | cudaMemcpy2DAsync |
| 0.07% | 2.7961ms | 45 | 62.134us | 9.4710us | 187.97us | cudaMalloc |
| 0.04% | 1.6533ms | 29 | 57.010us | 5.1150us | 650.77us | cudaStreamSynchronize |
| 0.03% | 1.3740ms | 4 | 343.51us | 341.96us | 344.80us | cuDeviceTotalMem |
| 0.02% | 1.0064ms | 352 | 2.8590us | 502ns | 77.755us | cuDeviceGetAttribute |
| 0.02% | 759.06us | 112 | 6.7770us | 866ns | 296.13us | cudaEventCreateWithFlags |
| 0.01% | 530.04us | 28 | 18.930us | 9.7590us | 55.998us | cudaLaunch |
| 0.01% | 463.17us | 6 | 77.195us | 26.242us | 129.89us | cudaMemcpy |
| 0.01% | 330.98us | 4 | 82.743us | 21.216us | 228.00us | cudaStreamCreate |
| 0.00% | 125.71us | 4 | 31.426us | 22.651us | 35.699us | cuDeviceGetName |
| 0.00% | 111.02us | 158 | 702ns | 520ns | 1.6090us | cudaSetupArgument |
| 0.00% | 99.803us | 104 | 959ns | 684ns | 2.5010us | cudaDeviceGetAttribute |
| 0.00% | 80.671us | 34 | 2.3720us | 801ns | 6.8400us | cudaSetDevice |


```

0.00% 45.615us    2 22.807us 20.577us 25.038us
cudaStreamCreateWithPriority
0.00% 36.588us   28 1.3060us    655ns 3.9660us cudaConfigureCall
0.00% 34.730us    1 34.730us 34.730us 34.730us cudaMemsetAsync
0.00% 30.635us   10 3.0630us 1.4260us 6.9790us cudaGetDevice
0.00% 15.498us   20    774ns    575ns 1.1170us cudaPeekAtLastError
0.00% 8.5720us    8 1.0710us    530ns 2.6530us cudaGetLastError
0.00% 7.9230us    1 7.9230us 7.9230us 7.9230us cudaEventQuery
0.00% 6.8350us    6 1.1390us    545ns 2.7790us cuDeviceGetCount
0.00% 5.0810us    6    846ns    687ns 1.1550us cuDeviceGet
0.00% 4.1240us    3 1.3740us 1.3230us 1.4450us cuInit
0.00% 3.9170us    2 1.9580us 1.6040us 2.3130us
cudaDeviceGetStreamPriorityRange
0.00% 2.8320us    1 2.8320us 2.8320us 2.8320us cudaEventRecord
0.00% 2.7000us    3    900ns    896ns    907ns cuDriverGetVersion
0.00% 1.8210us    1 1.8210us 1.8210us 1.8210us cudaGetDeviceCount

```

3. nvprof profiling the execution (10 data size):

```

* Running nvprof python m3.1.py 10
Loading fashion-mnist data...
done
Loading model...
==314== NVPROF is profiling process 314, command: python m3.1.py 10
done

```

```

New Inference
Op Time: 0.000260
Op Time: 0.000733

```

Correctness: 1.0 Model: ece408

```

==314== Profiling application: python m3.1.py 10

```

```

==314== Profiling result:

```

| Time (%) | Time | Calls | Name |
|----------|-----------|-------|--|
| 38.50 | 897.27 us | 2 | mxnet::op::forward_kernel |
| 36.27 | 845.40 us | 14 | [CUDA memcpy HtoD] |
| 10.71 | 249.60 us | 1 | void sgemm_largek_lds64 |
| 7.67 | 178.85 us | 14 | void mshadow::cuda::MapPlanKernel<...> |
| 1.74 | 40.448 us | 1 | sgemm_sm35_ldg_tn_32x16x64x8x16 |
| 1.37 | 31.840 us | 2 | void mshadow::cuda::MapPlanKernel |

| | | | |
|-------------|------------------|---|--|
| 1.11 | 25.823 us | 2 | void cudnn::detail::activation_fw_4d_kernel<...> |
| 0.81 | 18.976 us | 1 | void cudnn::detail::pooling_fw_4d_kernel<...> |
| 0.63 | 14.720 us | 1 | void mshadow::cuda::SoftmaxKernel<...> |
| 0.48 | 11.168 us | 1 | void mshadow::cuda::MapPlanKernel<...> |
| 0.27 | 6.2720 us | 2 | void mshadow::cuda::MapPlanKernel<...> |
| 0.20 | 4.6400 us | 1 | [CUDA memcpy DtoH] |
| 0.18 | 4.1280 us | 1 | void scal_kernel<...> |
| 0.06 | 1.4400 us | 1 | [CUDA memset] |

==314== API calls:

| Time(%) | Time | Call | Avg | Min | Max | Name |
|----------------------------------|----------|------|----------|----------|----------|---------------------------|
| 47.16% | 2.03746s | 16 | 127.34ms | 19.098us | 1.01841s | cudaStreamCreateWithFlags |
| 29.59% | 1.27843s | 10 | 127.84ms | 1.2330us | 363.76ms | cudaFree |
| 22.87% | 988.11ms | 27 | 36.597ms | 337.32us | 979.12ms | cudaMemGetInfo |
| 0.09% | 3.9655ms | 4 | 991.36us | 23.032us | 3.8549ms | cudaStreamCreate |
| 0.06% | 2.5601ms | 45 | 56.891us | 7.8580us | 208.19us | cudaMalloc |
| 0.06% | 2.4786ms | 9 | 275.40us | 12.252us | 758.74us | cudaMemcpy2DAsync |
| 0.03% | 1.3842ms | 4 | 346.06us | 340.94us | 360.30us | cuDeviceTotalMem |
| 0.02% | 1.0547ms | 352 | 2.9960us | 516ns | 117.54us | cuDeviceGetAttribute |
| 0.02% | 988.21us | 6 | 164.70us | 5.2810us | 702.96us | cudaDeviceSynchronize |
| 0.02% | 935.43us | 29 | 32.256us | 5.7630us | 218.72us | cudaStreamSynchronize |
| 0.02% | 859.79us | 112 | 7.6760us | 930ns | 294.14us | cudaEventCreateWithFlags |
| 0.01% | 538.01us | 28 | 19.214us | 9.2970us | 58.700us | cudaLaunch |
| 0.01% | 477.63us | 6 | 79.605us | 42.779us | 130.92us | cudaMemcpy |
| 0.00% | 120.71us | 4 | 30.176us | 18.506us | 35.825us | cuDeviceGetName |
| 0.00% | 109.93us | 158 | 695ns | 414ns | 1.7770us | cudaSetupArgument |
| 0.00% | 100.97us | 104 | 970ns | 571ns | 2.1010us | cudaDeviceGetAttribute |
| 0.00% | 83.217us | 34 | 2.4470us | 941ns | 7.5720us | cudaSetDevice |
| 0.00% | 42.570us | 2 | 21.285us | 20.207us | 22.363us | |
| cudaStreamCreateWithPriority | | | | | | |
| 0.00% | 37.688us | 28 | 1.3460us | 628ns | 4.0480us | cudaConfigureCall |
| 0.00% | 35.023us | 1 | 35.023us | 35.023us | 35.023us | cudaMemsetAsync |
| 0.00% | 31.994us | 10 | 3.1990us | 1.8510us | 7.2940us | cudaGetDevice |
| 0.00% | 16.583us | 20 | 829ns | 596ns | 1.1390us | cudaPeekAtLastError |
| 0.00% | 9.4350us | 8 | 1.1790us | 555ns | 3.8670us | cudaGetLastError |
| 0.00% | 6.7170us | 6 | 1.1190us | 525ns | 2.4450us | cuDeviceGetCount |
| 0.00% | 6.5240us | 1 | 6.5240us | 6.5240us | 6.5240us | cudaEventQuery |
| 0.00% | 5.8530us | 6 | 975ns | 673ns | 1.9630us | cuDeviceGet |
| 0.00% | 4.5010us | 3 | 1.5000us | 1.3600us | 1.6470us | cuInit |
| 0.00% | 3.8660us | 2 | 1.9330us | 1.5120us | 2.3540us | |
| cudaDeviceGetStreamPriorityRange | | | | | | |
| 0.00% | 2.8840us | 3 | 961ns | 956ns | 968ns | cuDriverGetVersion |
| 0.00% | 2.5270us | 1 | 2.5270us | 2.5270us | 2.5270us | cudaEventRecord |
| 0.00% | 1.6070us | 1 | 1.6070us | 1.6070us | 1.6070us | cudaGetDeviceCount |

4. Op Times Comparison (CPU vs GPU):

a). 10000 images (default):

| | CPU Time | GPU Time |
|------------------------------|-------------|------------|
| First Layer Op Time: | 6.60747 s | 0.355486 s |
| Second Layer Op Time: | 19.537141 s | 0.652564 s |

b). 100 images:

| | CPU Time | GPU Time |
|------------------------------|------------|------------|
| First Layer Op Time: | 0.065591 s | 0.003272 s |
| Second Layer Op Time: | 0.194542 s | 0.006654 s |

c). 10 images:

| | CPU Time | GPU Time |
|------------------------------|------------|------------|
| First Layer Op Time: | 0.006572 s | 0.000260 s |
| Second Layer Op Time: | 0.019520 s | 0.000733 s |

4. Discussion:

In this milestone we have completed a GPU implementation of the forward convolution. Though there is further optimization added into the new code, we could still observe the significant speed-boost brought by CPU computation. For the default data size, the time spent has been shrunk from 6.6 s to within half a second. This much performance enhancement is exactly why GPU computing is so powerful in many occasions.